Visualizing Billions of Data Points: Doing It Right

Alexey Zaytsev,

Skoltech, CDISE

18 January

Slides are based on http://go.continuum.io/visualizing-billions-data-points/

Overview

1. Visualizing big data — what is the problem?
2. Datashading
3. Datasets:
    - 10 million points of NYC Taxi data
    - 3 billion points of OpenStreetMap data
    - 300 million points of US Census data

Billions and billions…

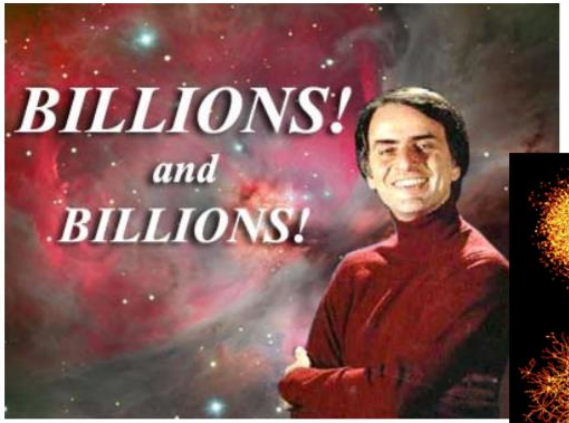

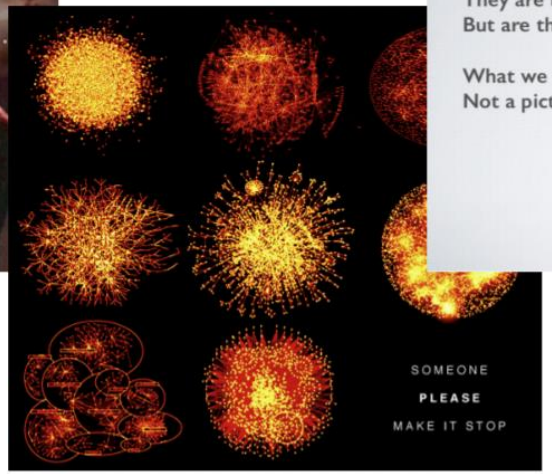

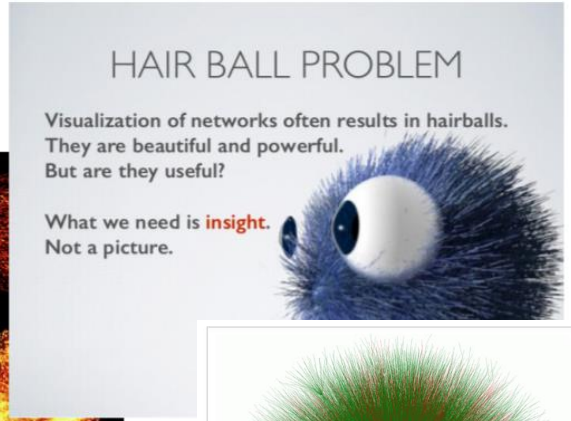

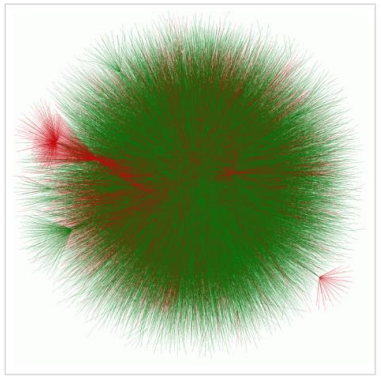

https://networkscience.wordpress.com/2016/06/22/no-hairball-the-graph-drawing-experiment/

Big data magnifies small problems

Big data presents storage and computation problems

More importantly, standard plotting tools have problems that are magnified by big data:

- Overdrawing/Overplotting
- Saturation
- Undersaturation
- Binning issues

We'll first explain these problems, and then present *datashading* technique to address them.

Overdrawning

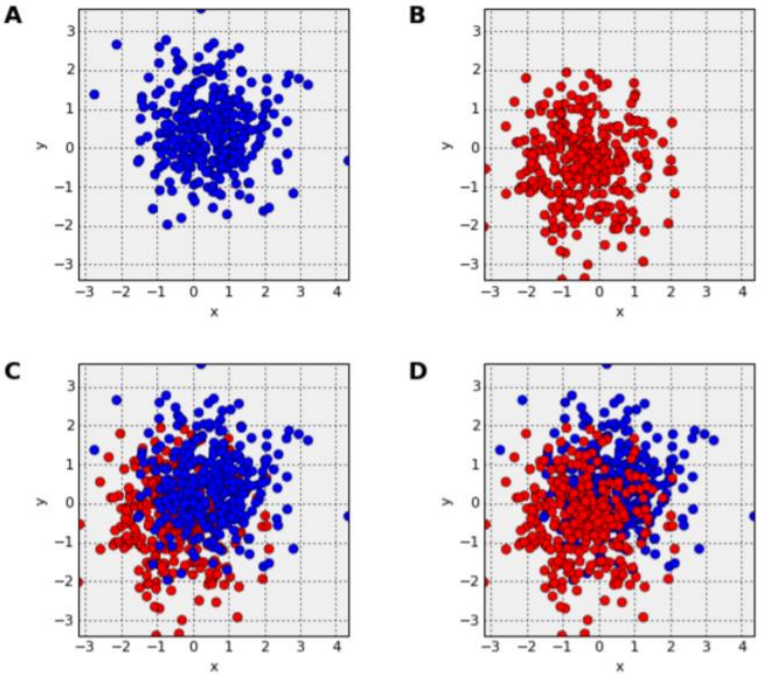

- For a scatterplot, the order in which points are drawn is very important
- The same distribution can look entirely different depending on plotting order
- Last data plotted overplots

Overdrawning

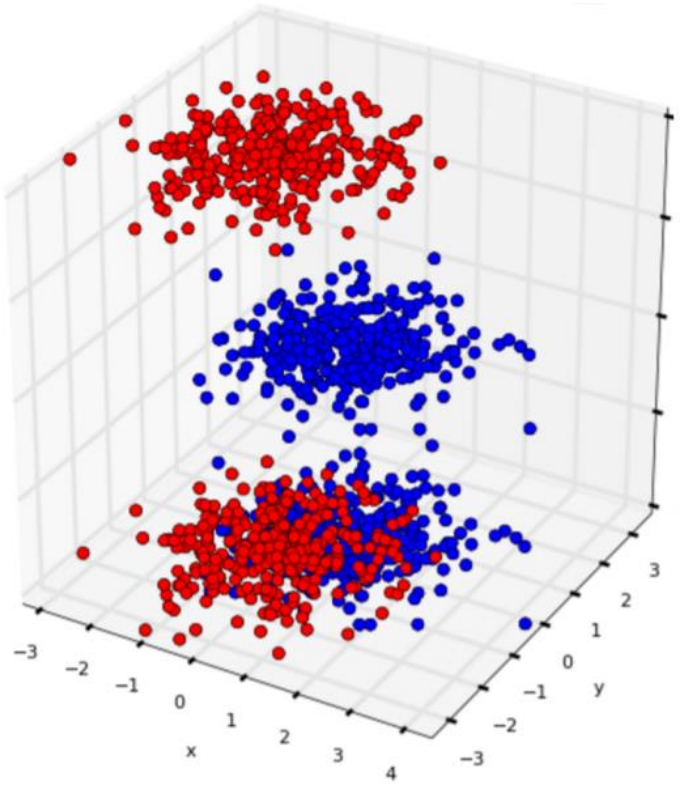

- Underlying issue is just occlusion
- Same problem happens with one category, but less obvious
- a Can prevent occlusion using transparency

Saturation

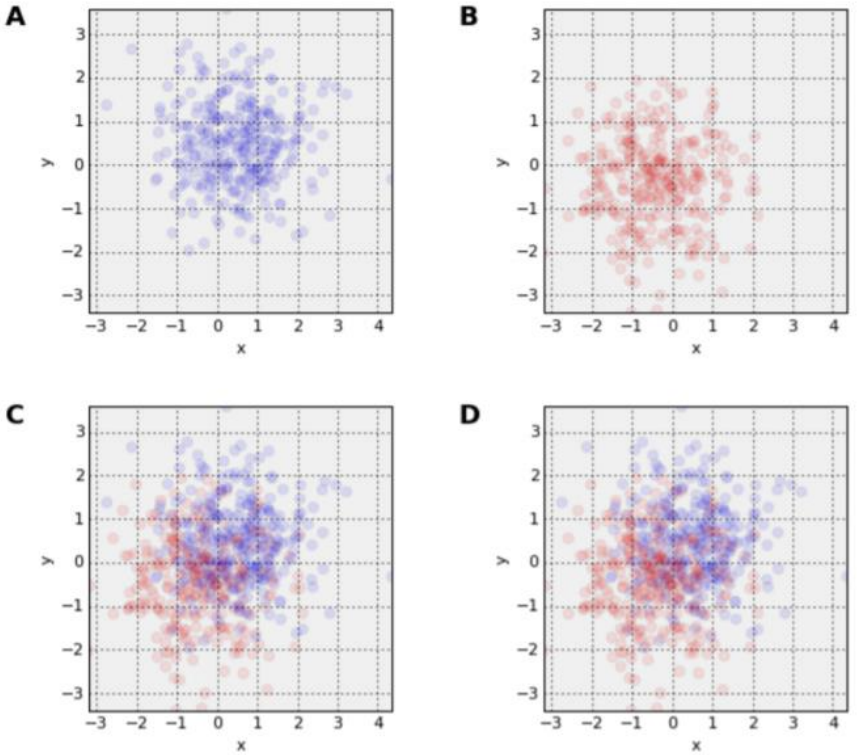

- For alpha = 0.1, up to 10 points can overlap before saturating the available brightness
- Now the order of plotting matters less
- After 10 points, first-plotted data still lost
- For one category, 10, 20, or 2000 points overlapping will look identical

Saturation

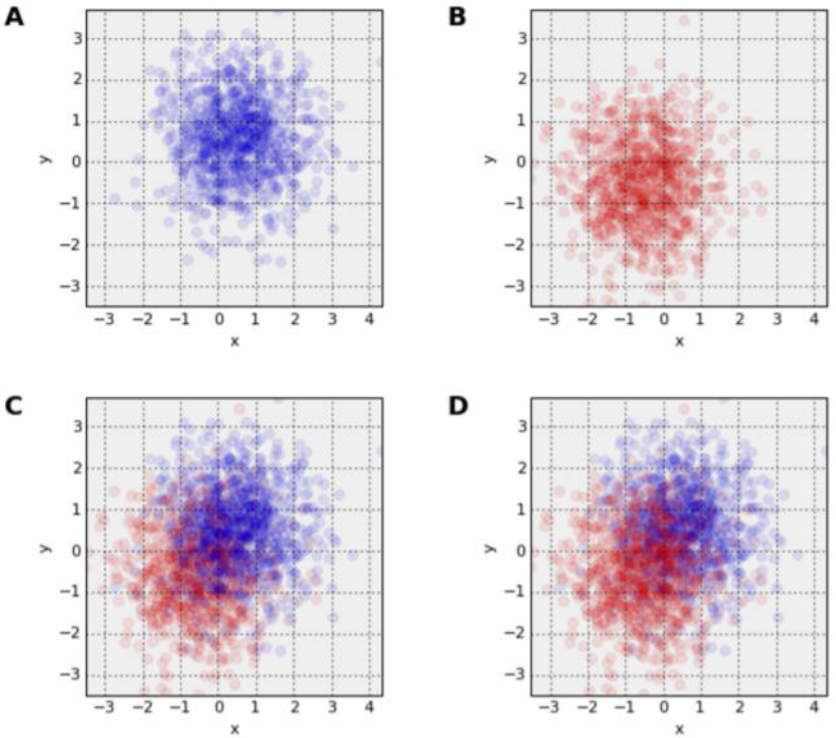

Same alpha value, more points:

- Now is highly misleading
- alpha value depends on size, overlap of dataset
- Difficult-to-set parameter, hard to know when data is misrepresented

Saturation

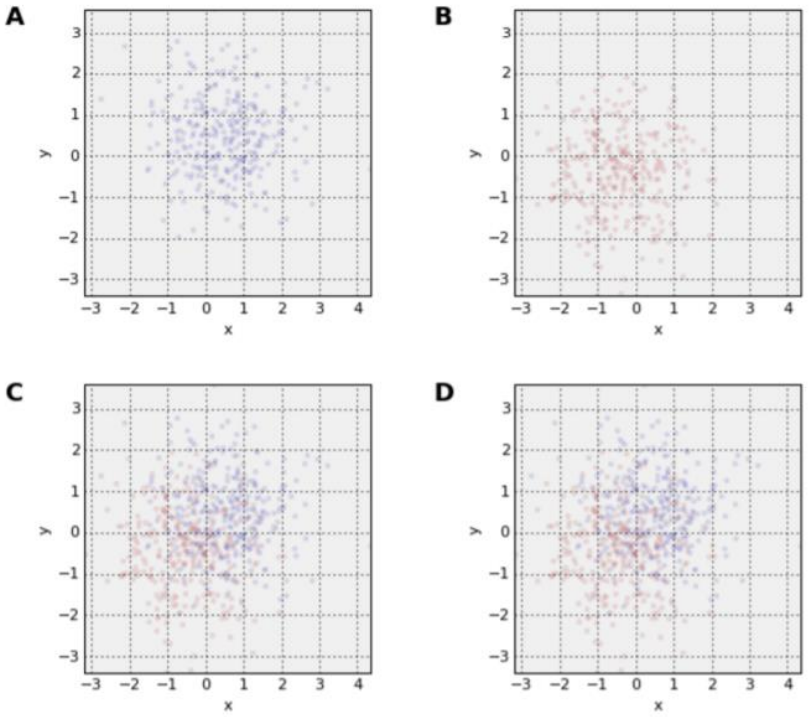

Can try to reduce point size to reduce overplotting and saturation

- Now points are hard to see, with no guarantee of avoiding problems
- Another difficult-to-set parameter
- For really big data, scatterplots start to become very inefficient: many points per pixel — may as well be binning by pixel

Binning

- Can use heatmap instead of scatter
- Avoids saturation by auto-ranging on bins
- Result independent of data size
- Here two merged normal distributions look very different at different binning
- Another difficult-to-set parameter

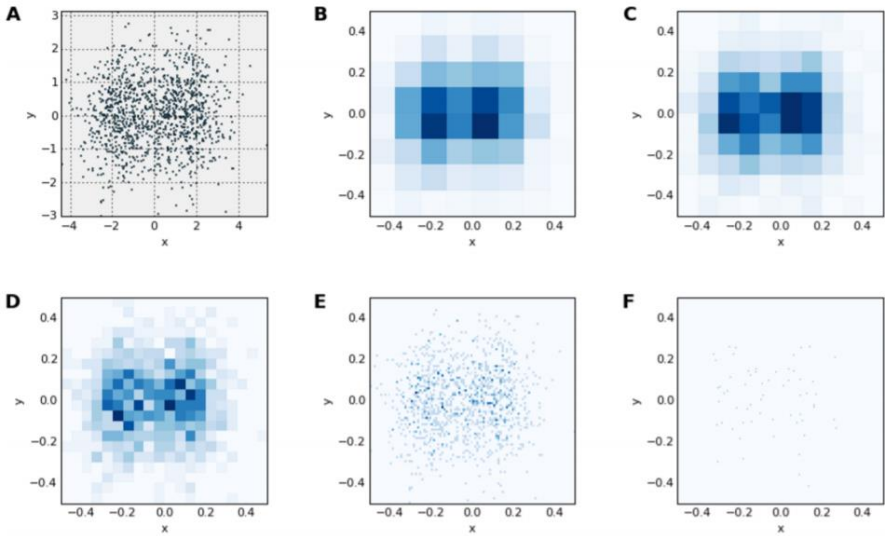

Main challenges while plotting big data

- When exploring really big data, the visualization is all you have — there's no way to look at each of the individual data points
- Common plotting problems can lead to completely incorrect conclusions based on misleading visualizations
- Slow processing makes trial and error approach ineffective

When data is large, you don't know when the plot is lying.

Datashading

- Flexible, configurable pipeline for automatic plotting
- Prevents overplotting, saturation, and undersaturation •
- Mitigates binning issues by interactivity, even of very large datasets on ordinary machines
- Allows rapid iteration of visual styles & configs, interactive selections and filtering, to support data exploration

Datashading pipeline starts with projection

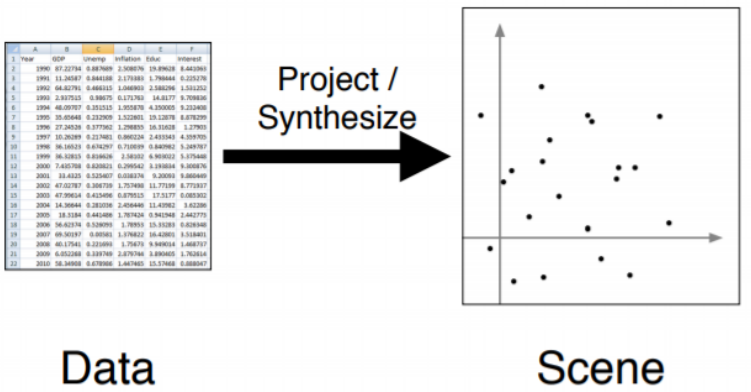

- Stage 1: select variables (columns) to project onto the screen
- Data often filtered at this stage

Datashading pipeline continues with aggregation

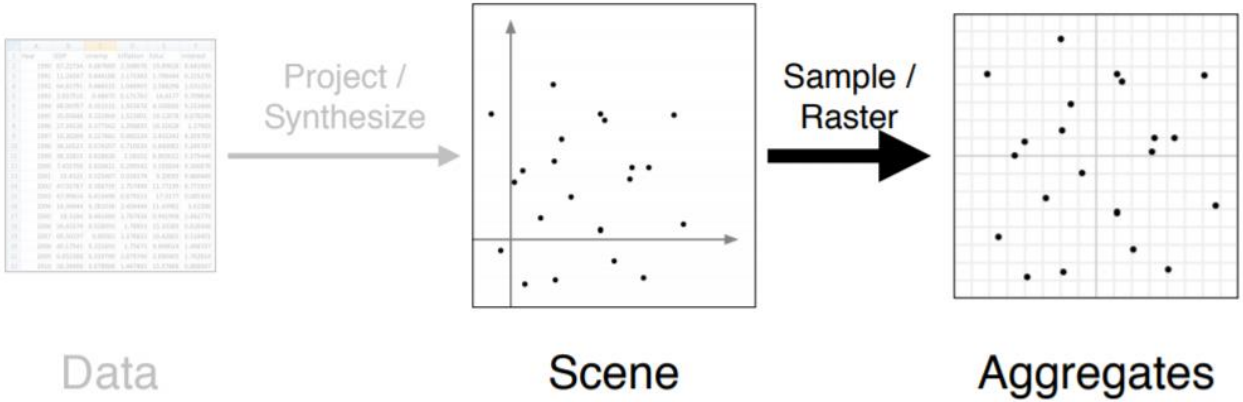

- Stage 2: Aggregate data into a fixed set of bins
- Each bin yields one or more scalars (total count, mean, stddev, etc.)

Datashading pipeline ends with transfer

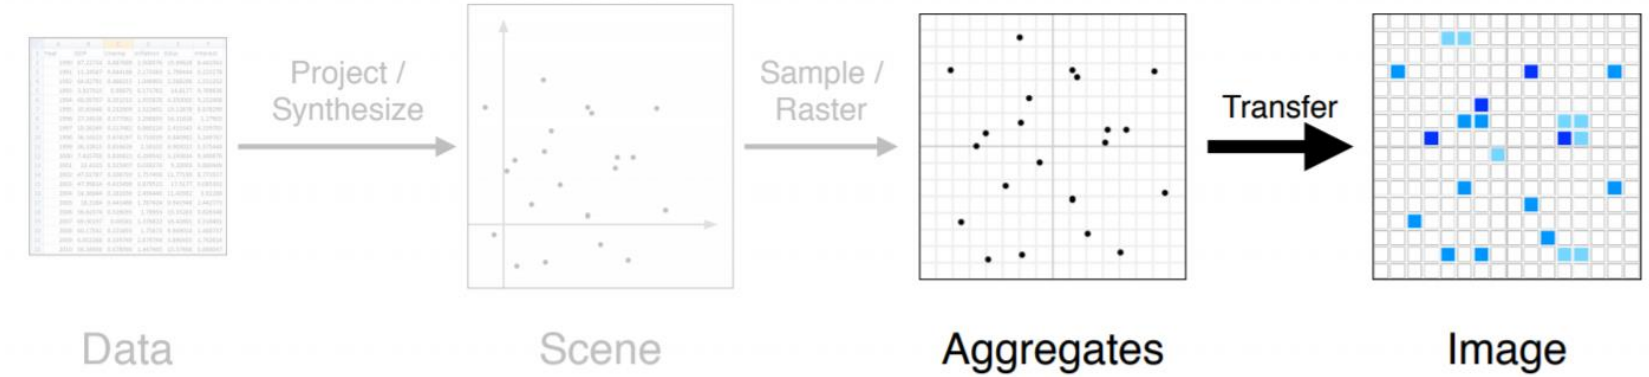

- Stage 3: Transform data using one or more transfer functions, culminating in a function that yields a visible image
- Each stage can be replaced and configured separately

NYC taxi data

- Data for 10 million New York City taxi trips
- Even 100,000 points gets slow for scatterplot
- Parameters usually need adjusting for every zoom
- True relationships within data not visible in std plot

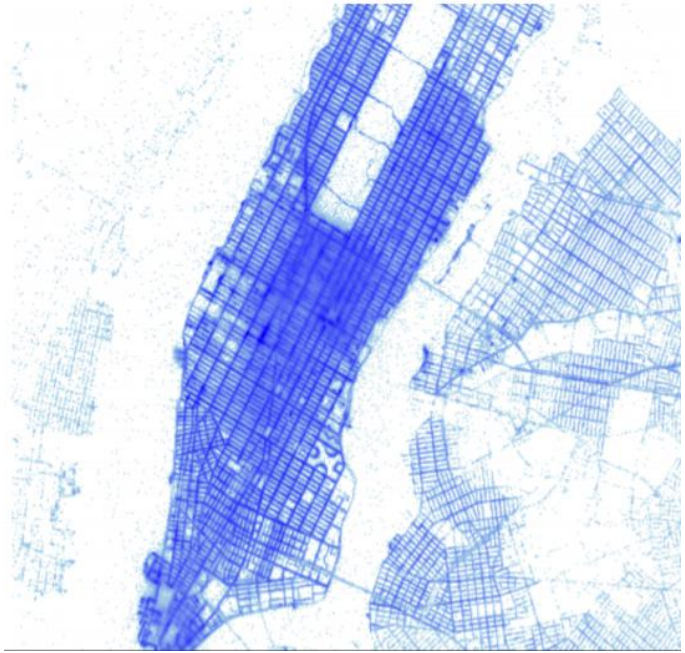

Datashading automatically reveals the entire dataset, including outliers, hot spots, and missing data

Billion OSM points

Open Street Map data:

- About 3 billion GPS coordinates
- https://blog.openstreetmap.org/2012/04/01/bulk-gps-point-data/
- This image was rendered in one minute on a standard MacBook with 16 GB RAM
- Renders in 7 seconds on a 128GB Amazon EC2 instance

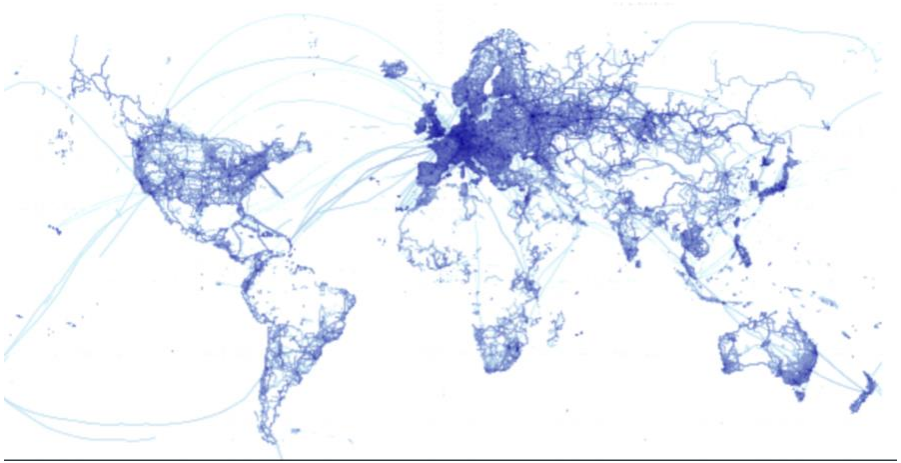

Billion OSM points

- OSM database is 40GB, larger than the 16GB RAM on this machine
- Fast out-of-core operation powered by: •
  - Numba: generates fast C and GPU code from Python source
  - Dask: Parallelizes tasks
- datashader source code is all Python

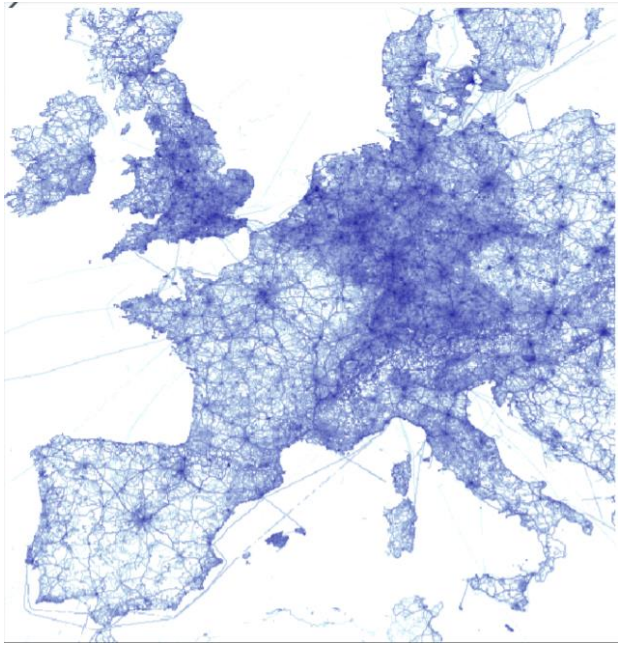

300 million points Census data

- One point per person
- 300 million total
- Categorized by race
- Datashading shows faithful distribution per pixel

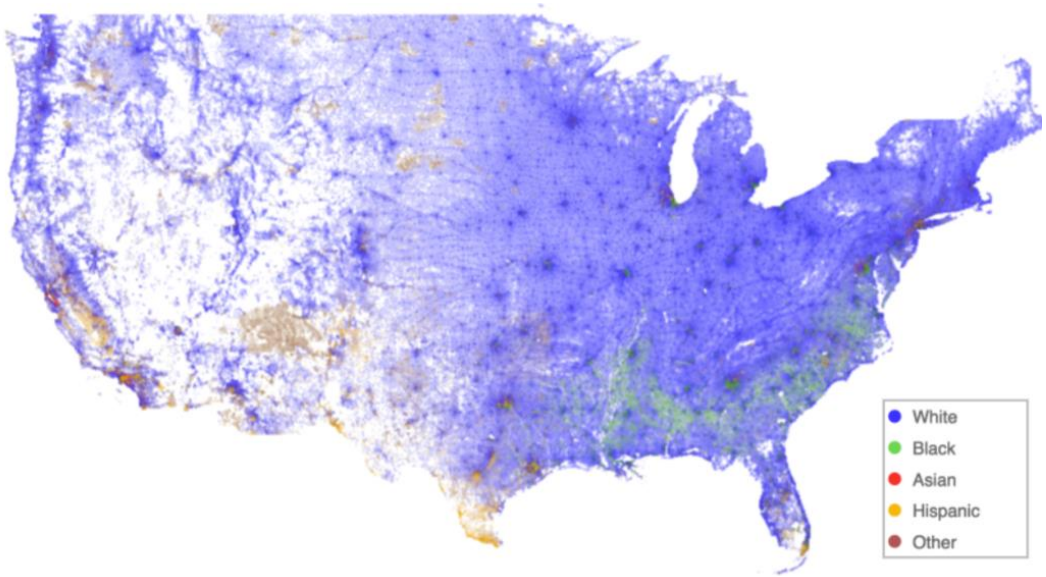

Example of data exploration with datashader

http://datashader.org/topics/nyc_taxi.html